Online Correction of Semantic Drift in Dense Retrieval via Implicit Feedback and Sparse Autoencoders

Anonymous authors
Paper under double-blind review

Keywords: Dense Retrieval, Semantic Drift, Reinforcement Learning, Sparse Autoencoders, RLHF, Zero-Shot

Summary

A novel approach to correcting semantic drift in dense vector retrieval systems by modeling user implicit feedback through a PPO agent bound by Sparse Autoencoders. The transition from lexical matching to dense vector representation has fundamentally altered information retrieval architectures. While dense retrieval captures deep contextual semantics, it introduces interpretability and control challenges, often leading to semantic drift. This paper proposes leveraging implicit user feedback as an online correction mechanism, formulating the retrieval correction process as an Implicit Reinforcement Learning from Human Feedback (RLHF) objective. By resolving the continuous-space intractability with a structural bottleneck derived from Sparse Autoencoders (SAEs), we train a highly lightweight Proximal Policy Optimization (PPO) agent to execute Negative Semantic Masking at query time, establishing a true zero-shot offline training protocol for domain adaptation on BEIR datasets with minimal computational overhead, empirically validating its topological fidelity and outperforming foundational baselines like Rocchio and HyDE.

Contribution(s)

1. We formulate semantic drift correction in dense retrieval as an MDP bounded by sparse dictionary features, explicitly solving the continuous-space intractability problem without requiring full-model contrastive retraining.
   Context: None
2. We introduce Context-Aware Latent-to-Dense Routing, demonstrating that implicit user feedback can guide a PPO agent to execute Negative Semantic Masking and Rank-Proportional Dynamic Torque at query time based solely on observable state dynamics (Retrieval Entropy and Facet Variance).
   Context: None
3. We establish a robust true zero-shot online correction protocol. By training offline on MS MARCO, we demonstrate out-of-the-box transfer on complex subsets like ArguAna, outperforming heuristic (Rocchio) and generative (HyDE) zero-shot baselines while rigidly preserving topological fidelity (Spearman $\rho = 0.9906$).
   Context: None

Online Correction of Semantic Drift in Dense Retrieval via Implicit Feedback and Sparse Autoencoders

Anonymous authors
Paper under double-blind review

Abstract

The transition from lexical matching to dense vector representation has fundamentally altered information retrieval architectures. While dense retrieval captures deep contextual semantics, it introduces interpretability and control challenges, often leading to semantic drift. This paper proposes leveraging implicit user feedback as an online correction mechanism, formulating the retrieval correction process as an Implicit Reinforcement Learning from Human Feedback (RLHF) objective. By resolving the continuous-space intractability with a structural bottleneck derived from Sparse Autoencoders (SAEs), we train a highly lightweight Proximal Policy Optimization (PPO) agent to execute Negative Semantic Masking at query time, establishing a true zero-shot offline training protocol for domain adaptation on BEIR datasets with minimal computational overhead, empirically validating its topological fidelity and outperforming foundational baselines like Rocchio and HyDE.

1 Introduction

The transition from lexical matching to dense vector representation has fundamentally altered information retrieval architectures (Karpukhin et al., 2020; Izacard et al., 2022). By projecting queries and documents into a shared continuous latent space, dense retrieval captures deep contextual semantics that term-frequency methods systematically miss. However, this architectural shift introduces severe interpretability and control challenges. High-dimensional embeddings inherently entangle distinct semantic properties, leading to representation collapse or semantic drift (Zhan et al., 2021). In these failure modes, models retrieve documents based on spurious distributional overlaps or dominant but irrelevant vector magnitudes rather than strict contextual relevance (Thakur et al., 2021). Correcting this drift traditionally necessitates computationally expensive offline fine-tuning over newly mined hard-negative distributions, a paradigm that is inherently reactive and scales poorly in dynamic retrieval environments.

To bypass the latency and computational overhead of offline retraining, we propose leveraging implicit user feedback as an online correction mechanism. In production retrieval systems, user interactions such as skipping top-ranked results to click on lower-ranked documents provide a natural, continuous reward signal (Joachims, 2002; Afsar et al., 2022). Formulating the retrieval correction process as an Implicit Reinforcement Learning from Human Feedback (RLHF) objective (Christiano et al., 2017) utilizing historical interaction logs as preference signals—offers a mathematically principled way to steer queries dynamically without requiring manual annotations. Yet, deploying reinforcement learning algorithms directly over entangled, high-dimensional dense spaces is computationally intractable and suffers from severe sample inefficiency.

We resolve this intractability by imposing a structural bottleneck derived from mechanistic interpretability, specifically utilizing Sparse Autoencoders (SAEs) (Cunningham et al., 2024; Elhage et al., 2022; Bricken et al., 2023). We propose a two-phase architecture. First, a Sparse State Encoder projects the dense query and the latent vocabulary of its top Pseudo-Relevance Feedback

(PRF) documents into an overcomplete, non-negative dictionary of isolated semantic facets, yielding a sparse activation vector $x \in \mathbb{R}^k$ where $k \gg d$. Second, we formalize the correction process as a Markov Decision Process (MDP) and train a highly lightweight Proximal Policy Optimization (PPO) agent (Schulman et al., 2017).

To overcome semantic blindness, the agent observes a Context-Aware State that concatenates the original dense query with the sparse PRF extraction. Driven by the implicit click-based reward signal, the PPO agent learns a steering policy $\pi_\theta(a|s)$ to execute Negative Semantic Masking, outputting a bi-directional steering matrix that actively subtracts spurious noise. Finally, to protect high-performing queries from exploration degradation, the projected dense shift is scaled by a Rank-Proportional Dynamic Torque before executing the inference-time correction.

2 Related Work

Our approach intersects three historically distinct domains: the application of reinforcement learning to information retrieval, the emerging field of mechanistic interpretability, and the architectural design of controllable sparse retrieval systems.

2.1 Reinforcement Learning in Information Retrieval

The formulation of search and recommendation as sequential decision-making processes has extensive precedent (Afsar et al., 2022). Early efforts successfully modeled session-based search and document ranking as Markov Decision Processes (MDPs), utilizing multi-armed bandits to balance exploration and exploitation in listwise ranking (Hofmann et al., 2013; Wei et al., 2017; Hu et al., 2018). Subsequent works extended reinforcement learning to query formulation, training discrete policy agents to sequentially append or modify search tokens based on terminal retrieval rewards (Nogueira & Cho, 2017; Buck et al., 2018). Parallel advances in recommendation systems heavily leveraged deep reinforcement learning to optimize long-term user engagement and page-wise recommendations (Ouyang et al., 2021; Zou et al., 2019; Zheng et al., 2018; Zhao et al., 2018), utilizing techniques such as off-policy correction (Chen et al., 2019) and dynamic reasoning over knowledge graphs (Yao et al., 2018).

More recently, the paradigm of Reinforcement Learning from Human Feedback (RLHF) (Christiano et al., 2017) has scaled to instruct-align large language models using proximal policy optimization against learned reward models (Ouyang et al., 2022; Stiennon et al., 2020; Bai et al., 2022; Ziegler et al., 2019). However, a critical limitation pervades the application of RL to retrieval: existing agents fundamentally operate on discrete token spaces (e.g., query expansion) or perform discrete combinatorics over candidate document lists (e.g., re-ranking). Attempting to train a continuous-action PPO policy directly on an entangled dense latent space $\mathbb{R}^d$ without an orthogonal, human-interpretable basis yields highly unstable gradients and severe sample inefficiency. Continuous vector perturbation in dense spaces fails to cleanly isolate semantic features, resulting in out-of-distribution catastrophic failures.

2.2 Mechanistic Interpretability and Sparse Autoencoders

The structural opacity of dense embeddings stems primarily from the phenomenon of superposition, wherein neural networks represent more features than they have dimensions by packing them into almost-orthogonal latent directions (Elhage et al., 2022). Foundational work in mechanistic interpretability initially sought to decode this polysemanticity by analyzing individual neurons and multimodal circuits (Olah et al., 2020; Goh et al., 2021), or by probing Transformer feed-forward layers as explicit key-value conceptual memories (Yun et al., 2020; Geva et al., 2022; Meng et al., 2022). However, probing raw activations remains insufficient for systematic intervention due to the inherent entanglement of the latent bases (Lee et al., 2023; Arora et al., 2018).

Recent advances have demonstrated that Sparse Autoencoders (SAEs) can effectively untangle these representations. By training a shallow autoencoder with an L_1 sparsity penalty on a wide bottleneck, SAEs project dense activations into a high-dimensional, overcomplete dictionary of monosemantic features (Cunningham et al., 2024; Bricken et al., 2023; Templeton et al., 2024; Gao et al., 2024). This allows for the precise isolation of semantic concepts, facilitating automated circuit discovery and the tracking of representational grokking (Conmy et al., 2023; Nanda et al., 2023; Wang et al., 2023). In this work, we treat the SAE not as an analytical probe, but as a rigid structural tool: it serves as the requisite state-space transformation, linearizing the dense embedding $v \in \mathbb{R}^d$ into a discrete, manipulable sparse vector $x \in \mathbb{R}^k$. This transforms the intractable continuous latent steering problem into a highly constrained, multi-discrete MDP suitable for lightweight RL.

2.3 Controllable and Sparse Retrieval

The empirical success of dual-encoder dense retrieval (Karpukhin et al., 2020; Izacard et al., 2022) is often compromised by the system's susceptibility to semantic drift and out-of-domain degradation (Zhan et al., 2021). Traditionally, correcting the retrieval manifold requires computationally expensive offline retraining via Approximate Nearest Neighbor (ANN) hard-negative mining (Qu et al., 2021; Hofstätter et al., 2021; Xiong et al., 2021). These contrastive updates are reactive and fail to offer inference-time steerability.

To regain explicit control, hybrid and sparse retrieval paradigms have re-emerged. Traditional lexical baselines like BM25 (Robertson et al., 1994) have been augmented by neural term-weighting models such as docTquery and DeepCT (Nogueira & Lin, 2019; Dai & Callan, 2019). More advanced learned sparse models, notably SPLADE and CoIL, project document semantics directly into the exact vocabulary space, generating interpretable, human-readable sparse vectors (Formal et al., 2021a;b; Gao et al., 2021). Alternatively, late-interaction architectures like ColBERT preserve fine-grained token-level similarities while delaying the representation bottleneck (Khattab & Zaharia, 2020; Santhanam et al., 2022; Lin et al., 2021; Luan et al., 2021). While these models successfully recover exact-match signals and offer term-level interpretability, they are rigidly bound to the discrete lexical vocabulary. They cannot perform continuous, granular amplification or suppression of abstract, multi-token semantic concepts. By coupling an SAE with an RL agent, our framework bridges this gap, providing continuous semantic steerability without the need for offline hard-negative retraining.

Recent works exploring SAEs for dense retrieval (e.g., CL-SR and *k*-sparse SAEs) similarly demonstrate that modifying latent features can steer retrieval. While these works utilize direct latent manipulation and contrastive objectives, our RL-driven approach complements them by providing a dynamic, policy-based routing mechanism. Furthermore, in zero-shot domain adaptation, methods like MoDIR explicitly target representation invariance, while advanced dense models (e.g., E5, BGE) offer strong uncorrected limits. We leave direct empirical comparisons to these massive architectures to future work, constrained here by the explicit goal of demonstrating lightweight scalar routing over a frozen baseline.

3 Problem Formulation

To formally establish the intractability of direct continuous-space query correction and motivate our sparse Reinforcement Learning from Human Feedback (RLHF) framework, we first mathematically define the retrieval objective, the mechanics of semantic drift, and the fundamental limitations imposed by high-dimensional geometry.

3.1 Dense Retrieval as Maximum Inner Product Search

Standard dense information retrieval systems operate by embedding a user query q and a set of corpus documents $d \in \mathcal{C}$ into a shared, d-dimensional continuous latent space. Let $E_Q(\cdot)$ and $E_D(\cdot)$ denote the query and document encoders, respectively. The relevance score between a query

and a document is formulated as a Maximum Inner Product Search (MIPS) (Johnson et al., 2019; Shrivastava & Li, 2014):

$$S(q, d) = E_Q(q) \cdot E_D(d) \tag{1}$$

The retrieval objective is to retrieve the top-k documents that maximize this inner product. While highly efficient in production environments leveraging Approximate Nearest Neighbor (ANN) indices, this formulation implicitly assumes that the geometric proximity in $\mathbb{R}^d$ perfectly aligns with human semantic judgments.

3.2 Semantic Drift and the Dimensionality Bottleneck

A retrieval failure, or "semantic drift," occurs when the geometry of the latent space misaligns with true contextual relevance. Formally, a failure state exists when an irrelevant document $d_{irrelevant}$ scores higher than the contextually optimal target document d_{target}:

$$S(q, d_{irrelevant}) > S(q, d_{target}) \tag{2}$$

To correct this failure at inference time, one must theoretically apply a perturbation vector $\Delta_q \in \mathbb{R}^d$ such that the updated query representation $E'_q = E_Q(q) + \Delta_q$ reverses the inequality. However, solving for Δ_q directly in the original continuous latent space (e.g., $d = 768$) is mathematically ill-posed due to the curse of dimensionality (Beyer et al., 1999). In high-dimensional spaces, the distance contrast between the nearest and furthest neighbors approaches zero, meaning arbitrary continuous shifts disproportionately impact the global nearest-neighbor topology.

Furthermore, dense representations natively exhibit severe anisotropy—they are constrained to narrow cones within the vector space rather than being uniformly distributed (Ethayarajh, 2019). Because the latent basis vectors are highly entangled and polysemantic, applying a continuous spatial adjustment Δ_q to suppress a specific irrelevant concept simultaneously destroys unrelated semantic features. Consequently, naive continuous steering degrades the vector topology, necessitating the traditional, expensive fallback of offline whole-model contrastive retraining.

3.3 Episodic Formulation for RLHF

To bypass offline fine-tuning, we formulate the dynamic correction of $E_Q(q)$ as an episodic Markov Decision Process (MDP) (Sutton & Barto, 2018), defined by the tuple $\mathcal{M} = (\mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma)$. In this framework, search is inherently interactive. The environment encompasses the retrieval index and the user, while the agent is the query-steering mechanism.

State Space ($\mathcal{S}$): To overcome the "Anonymous Feature Trap," the state cannot simply be the query representation. We define a Context-Aware State $S \in \mathbb{R}^{898}$, concatenating the lossless 768-D dense query $E_Q(q)$ with a 128-D magnitude vector representing the semantic features of the local Pseudo-Relevance Feedback (PRF) neighborhood, alongside continuous rank-tracking scalars.

Action Space ($\mathcal{A}$): The agent outputs a continuous action vector $a_t \in [-2.0, 2.0]^{128}$. This bi-directional framing enables "Negative Semantic Masking," empowering the agent to explicitly subtract localized noise from the PRF extraction before the latent-to-dense projection.

Transition Dynamics ($\mathcal{P}$): The environment transitions deterministically. The agent's sparse action matrix is projected back into $\mathbb{R}^{768}$ using the Sparse Autoencoder's decoder weights. This dense shift is scaled by a rank-proportional dynamic torque, updating $E_Q(q)$ and yielding a newly sorted document list.

Reward Function ($\mathcal{R}$): The system receives a delayed, implicit reward based on simulated user interaction (Joachims, 2002). We formulate a continuous Deep-Rescue Logarithmic Reward proportional to the step-wise rank improvement of the target document, augmented with discrete positive spikes for breaching the Top 10 and Top 3 retrieval thresholds.

A fundamental challenge arises if we attempt to define the action space $\mathcal{A}$ as continuous perturbations in $\mathbb{R}^d$. Standard policy gradient methods fail to converge when the continuous action space

lacks an orthogonal, interpretable basis, as the agent cannot isolate specific semantic features to maximize the click-based reward. Resolving this intractability requires transforming the dense state space into a discrete, disentangled manifold prior to RL intervention.

By projecting the entangled $\mathbb{R}^{768}$ space into an interpretable $\mathbb{R}^{4096}$ sparse basis, we provide the RLHF agent with the necessary orthogonal controls to surgically alter semantic concepts. The agent learns to dynamically route queries at inference time, bridging the gap between exact lexical matching and dense conceptual understanding without requiring the costly architectural overhead of late-interaction re-ranking.

4 Methodology: RL-Driven Semantic Steering

To resolve the intractability of operating over a continuous, highly entangled dense space, we propose a two-phase architecture. We first project the retrieval representations into a sparse, linearizable state space, and subsequently define a highly constrained Markov Decision Process (MDP) optimized via Proximal Policy Optimization (PPO).

4.1 The Sparse State Encoder

The foundational layer of our methodology is the Sparse State Encoder, parameterized as a shallow autoencoder with a strict sparsity constraint. Let $E \in \mathbb{R}^d$ represent a dense query or document embedding extracted from a frozen dual-encoder retrieval model. We project this dense vector into a high-dimensional sparse space $F \in \mathbb{R}^k$ (where $k \gg d$) to isolate distinct semantic facets. The encoder mapping is defined as:

$$F = \text{ReLU}(W_{enc}E + b_{enc}) \tag{3}$$

where $W_{enc} \in \mathbb{R}^{k \times d}$ is the encoding projection matrix, $b_{enc} \in \mathbb{R}^k$ is the bias vector, and the ReLU activation ensures non-negativity, analogous to the biological plausibility of sparse coding (Olshausen & Field, 1996; Glorot et al., 2011). The decoder reconstructs the dense embedding as $\hat{E} = W_{dec}F + b_{dec}$.

To guarantee that the sparse features correctly encapsulate the semantic properties of the original continuous space, we optimize the autoencoder using a composite loss function combining Mean Squared Error for reconstruction fidelity and an L1 penalty for sparsity:

$$L_{SAE} = L_{MSE}(E, \hat{E}) + \lambda \|F\|_1 \tag{4}$$

where $\lambda = 1e-4$. This absolute sparsity constraint ensures that only a minimal, interpretable subset of semantic facets activates for any given query.

$$L_{IP} = \sum_{i,j \in \mathcal{B}} ((E_i \cdot E_j) - (\hat{E}_i \cdot \hat{E}_j))^2 \tag{5}$$

where $\mathcal{B}$ is a batch of embeddings. Preserving the inner product is a strictly enforced constraint; it guarantees that the Maximum Inner Product Search (MIPS) topology remains mathematically identical to the original dense environment, preventing structural collapse during transformation.

4.2 The Markov Decision Process (MDP)

Operating RL directly on the 4096-dimensional output of the SAE would reintroduce the curse of dimensionality, rendering gradient estimation highly inefficient (Lillicrap et al., 2016). Therefore, we formulate a highly compressed, dynamic MDP designed specifically for the topological constraints of Information Retrieval.

Contextual Vocabulary Extraction via PRF: Before defining the state space, the environment must extract the domain-specific vocabulary relevant to the user's query. We achieve this by projecting

the top $k = 10$ retrieved documents from the frozen dense baseline through the Sparse Autoencoder to retrieve their sparse representations, F_{d_i}. We construct a localized expansion candidate vector, F_{exp}, by aggregating the query's sparse representation F_q with a rank-weighted decay of the PRF neighborhood:

$$F_{exp} = F_q + \sum_{i=1}^{k} w_i F_{d_i} \tag{6}$$

where the weights w_i are linearly decayed from 1.0 to 0.1 based on their baseline MIPS rank, and subsequently normalized such that $\sum w_i = 1$. The environment then identifies the indices of the top 128 highest magnitude dimensions within F_{exp}. These 128 indices represent the most structurally significant semantic facets of the localized retrieval space, and they dynamically define the active operational bounds for the agent's state and action spaces.

State Space ($\mathcal{S}$): To avoid semantic blindness, the agent must understand both the original query intent and the context of its current retrieval neighborhood. We utilize a Context-Aware State $s_t \in \mathbb{R}^{898}$. First, the environment identifies expansion candidates by extracting the sparse representations of the top 10 documents retrieved by the base dense query, applying a rank-weighted decay (1.0 to 0.1). The state is the concatenation of: (1) the lossless 768-D dense query vector E_q, (2) the scalar values of the top 128 most active facets from the PRF neighborhood, (3) Retrieval Entropy (the softmax entropy of the top-k PRF distances), and (4) Facet Variance (variance of the active expansion candidates), ensuring no oracle metrics are utilized during inference.

Action Space ($\mathcal{A}$): The agent outputs a continuous vector of delta modifiers $\Delta M \in \mathbb{R}^{128}$ corresponding to the active PRF facets. To allow for "Negative Semantic Masking," we rigidly bound the active multipliers using a bi-directional clipping guardrail:

$$M_{t+1} \in [-2.0, 2.0] \tag{7}$$

This allows the agent to explicitly subtract spurious semantic clumps. The sparse matrix M is then projected back into the continuous space via the SAE decoder weights to form a dense shift vector. Crucially, to protect top-tier queries from PPO exploration noise, we scale this shift using Rank-Proportional Dynamic Torque. To define this mathematically, let H_{prf} represent the Retrieval Entropy of the baseline PRF documents, and H_{max} the maximum possible entropy. The dynamic torque multiplier τ is computed as a normalized scaling of this entropy constraint, reserving maximal leverage for queries with high semantic uncertainty:

$$\tau = \max\left(0.1, \min\left(1.0, \frac{H_{prf}}{H_{max}}\right)\right) \tag{8}$$

The final continuous dense query vector executed against the Approximate Nearest Neighbor (ANN) index is thus computed as:

$$E'_q = E_q + \tau(\Delta M W_{dec}^T) \tag{9}$$

where ΔM is the bi-directional sparse action matrix and W_{dec}^T is the transposed decoder weight matrix of the SAE. This formulation guarantees that queries already succeeding (low ambient entropy) receive near-zero physical leverage ($\tau \approx 0.1$), safeguarding them from PPO exploration noise, while uncertain queries receive maximal corrective torque.

The torque applied to the dense shift is clamped between 0.1 (for Rank 1 queries) and 1.0 (for deep failures), ensuring the agent only applies maximum physical leverage when a query requires rescuing.

Reward Function ($\mathcal{R}$): During offline training on MS MARCO, we shape simulated click feedback (derived from ground-truth relevance labels) into a continuous Deep-Rescue Logarithmic Reward (a shaped log-rank reward) to accelerate convergence prior to zero-shot deployment. The base reward is proportional to the logarithmic delta between the previous and current rank, preventing the agent from ignoring deeply buried relevant documents:

$$R_{dense} = 10.0 \times (\log_2(\text{Rank}_{old} + 1.0) - \log_2(\text{Rank}_{new} + 1.0)) \tag{10}$$

To incentivize terminal success, we apply massive discrete bonuses when the target document breaches critical visibility thresholds: $+20.0$ for entering the Top 10, and an additional $+10.0$ for breaching the Top 3.

4.3 Proximal Policy Optimization (PPO)

We optimize the steering policy using Proximal Policy Optimization (PPO) (Schulman et al., 2017), an actor-critic algorithm (Konda & Tsitsiklis, 1999) selected specifically for its gradient stability. Standard policy gradient algorithms, such as REINFORCE (Williams, 1992), exhibit high variance and allow unbounded step sizes that would irrecoverably destroy the precise query vector geometry.

The architecture comprises two separate Multi-Layer Perceptrons (MLPs): the Actor $\pi_\theta(a_t|s_t)$ which proposes the continuous steering vector ΔM, and the Critic $V_\phi(s_t)$ which estimates the expected return to minimize variance via Generalized Advantage Estimation (GAE) (Schulman et al., 2016), generating the advantage estimate $\hat{A}_t$.

To prevent destructive updates to the query representation, PPO constrains the policy shift at each training step. Let $r_t(\theta) = \frac{\pi_\theta(a_t|s_t)}{\pi_{\theta_{old}}(a_t|s_t)}$ denote the probability ratio between the updated and previous policies. The Actor is optimized using the clipped surrogate objective:

$$L^{CLIP}(\theta) = \hat{\mathbb{E}}_t \left[\min(r_t(\theta)\hat{A}_t, \text{clip}(r_t(\theta), 1-\epsilon, 1+\epsilon)\hat{A}_t) \right] \tag{11}$$

where $\epsilon = 0.2$ is the clipping hyperparameter. This objective restricts the magnitude of ΔM updates, ensuring the agent steadily and safely learns to penalize semantic drift without destabilizing the underlying dense manifold (Engstrom et al., 2020).

5 Experimental Setup

To rigorously evaluate the efficacy of our RL-driven semantic steering framework, we benchmark its performance across diverse, out-of-domain retrieval tasks. Our experimental design strictly constraints computational overhead, demonstrating that effective inference-time correction can be achieved without the massive hardware footprint typically required for standard hard-negative fine-tuning. All code, data splits, and model checkpoints required to reproduce our experiments are publicly available at https://anonymous.4open.science/r/IR-interpretability-RL-0632/.

5.1 Datasets

We evaluate our methodology utilizing five heterogeneous subsets from the BEIR (Benchmarking IR) zero-shot evaluation suite (Thakur et al., 2021). Because semantic drift predominantly occurs when models are deployed in domains drastically different from their pre-training distributions, we selected datasets spanning specialized scientific, financial, and conversational domains.

The selected datasets are detailed in Table 1. SciDocs (Cohan et al., 2020) evaluates direct citation prediction in scientific literature. ArguAna (Wachsmuth et al., 2018) tests the model's ability to retrieve highly abstract counter-arguments. NFCorpus (Boteva et al., 2016) provides a highly specialized medical vocabulary challenge. FiQA (Maia et al., 2018) evaluates financial question answering, which heavily relies on precise numerical and domain-specific semantic mappings. Finally, TREC-COVID (Voorhees et al., 2021) serves as a standard pandemic-era biomedical retrieval task. It is important to note for these latter two datasets that financial QA and pandemic nomenclature rely heavily on exact, highly specific jargon. While baseline sparse models undergo massive, domain-specific Masked Language Modeling (MLM) pre-training to memorize these tokens, our architecture's Sparse Autoencoder (SAE) is constrained to a lightweight 100-epoch training phase, providing a rigorous test of zero-shot domain adaptation under constrained vocabulary conditions.

Table 1: Benchmark Dataset Statistics

Dataset Name	Domain	Number of Queries Used	Number of Documents Used
SciDocs	Scientific Literature	800	12,502
ArguAna	Argument Retrieval	800	8,674
NFCorpus	Medical / Health	323	3,414
FiQA	Financial QA	648	14,549
TREC-COVID	Biomedical / News	50	17,537

**Note:* Reported latency reflects isolated routing/forward-pass overhead, not end-to-end ANN fetch time.

5.2 Baselines

We benchmark our RL-steered architecture against three distinct paradigms of neural information retrieval to isolate the comparative advantage of dynamic sparse correction:

Contriever (Dense Baseline) (Izacard et al., 2022): An unsupervised dense information retrieval model trained via contrastive learning. We utilize Contriever as the foundational, frozen base model for our architecture. The query and document representations $E \in \mathbb{R}^{768}$ generated by this model serve as the input to our Sparse Autoencoder.

Rocchio Vector PRF: A classical heuristic adaptation to dense space, where pseudo-relevance feedback documents are averaged and added to the query vector to shift it toward the target distribution.

HyDE (Hypothetical Document Embeddings): A zero-shot prompt-based method utilizing an instruction-tuned model (*flan-t5-small*) to generate a hypothetical relevant document, which is then embedded to perform the dense search.

ConSharp/SPLADE/ColBERT: For remaining datasets, we leave historical comparisons to exact-match sparse methods and high-latency late-interaction architectures as upper bounds for domain-specific fine-tuning.

5.3 Sparse Autoencoder Pre-Training & Topological Fidelity

Prior to RL optimization, we map the 768-D dense embeddings into an interpretable 4096-D latent space. To ensure strict structural zero-shot evaluation, the Sparse Autoencoder and the PPO agent are trained *offline* entirely on a disjoint subset of 10,000 queries from the MS MARCO dataset. The Sparse Autoencoder is trained for 100 epochs using the Adam optimizer (learning rate = 2e-3) and a batch size of 1024. The loss function combines Mean Squared Error (MSE) for reconstruction fidelity with an L1 sparsity penalty ($\lambda = 1e - 4$) to enforce monosemantic facet activation.

Before executing RL steering on the target domains, we computationally prove that this SAE acts as a safe, non-destructive structural bottleneck. By evaluating the Spearman Rank Correlation between the original Contriever dense MIPS rankings and the SAE-reconstructed MIPS rankings on the ArguAna dataset, we achieve a correlation of $\rho = 0.9906$. This confirms our latent projection preserves the fundamental metric geometry while introducing orthogonal sparsity.

5.4 Vectorized MDP Formulation & State Space

To ensure exact reproducibility and eliminate training bottlenecks, the environment is formulated as a fully vectorized Markov Decision Process (MDP) that processes all queries simultaneously using batched tensor operations (scatter_add_, torch.matmul).

For each query, the environment computes a Context-Aware State vector $S \in \mathbb{R}^{898}$. This observation space is a concatenation of:

- The lossless 768-D dense query context.
- A 128-D vector representing the magnitudes of the top expansion candidates extracted via Pseudo-Relevance Feedback (PRF). The PRF extracts the top 10 documents using a rank-weighted decay (1.0 to 0.1) to ensure the agent receives rich, domain-specific vocabulary.

- A scalar representing the $\log_2$ normalized current rank of the target document.
- A scalar representing the $\log_2$ step-wise delta in rank performance.

5.5 Proximal Policy Optimization (PPO) Setup

The steering policy is governed by a PPO agent utilizing a 3-layer Multi-Layer Perceptron (MLP) with Tanh activations for both the Actor and Critic networks (hidden dimension = 256). The agent maps the 898-D state space into a 128-D continuous action space.

Crucial constraints applied to the agent include:

- **Bi-Directional Steering:** Actions are clamped to $[-2.0, 2.0]$. This mathematical floor enables "Negative Semantic Masking" (i.e., using bounded signed scalar multipliers), allowing the agent to subtract false-positive noise explicitly from the dense vector.
- **Rank-Proportional Dynamic Torque:** To prevent PPO exploration noise from degrading queries that are already succeeding, the leverage applied to the dense shift is dynamically scaled based on the query's initial rank penalty. Torque is clamped between 0.1 (for perfect Rank 1 queries) and 1.0 (for deeply ranked failures).
- **Deep-Rescue Logarithmic Reward:** The agent receives a continuous reward proportional to the logarithmic delta between the previous and current rank, augmented with massive discrete spikes (+20.0) for breaching the NDCG@10 threshold, and (+10.0) for reaching the Top 3.

Reward Function Formulation: To formalize the "Deep-Rescue Logarithmic Reward," the episodic reward R_t at step t is calculated dynamically based on the reciprocal rank shift. Let ρ_{t-1} represent the query's document rank prior to the RL step, and ρ_t represent the rank post-intervention. The base reward is defined as $R_t = 10.0 \times (\log_2(\rho_{t-1} + 1.0) - \log_2(\rho_t + 1.0))$. This logarithmic scaling ensures that pushing a document from rank 1000 to 500 yields a comparable gradient update to pushing a document from rank 10 to 5, preventing the agent from ignoring deeply buried relevant documents.

Hyperparameter Optimization: The policy and value functions were optimized over 300 episodes using the Adam optimizer with a learning rate of $\eta = 1e-3$. To manage the tradeoff between exploration and exploitation in the continuous action space, we utilized Generalized Advantage Estimation (GAE) with a discount factor $\gamma = 0.99$ and a trace decay parameter $\lambda = 0.95$. The PPO clipping parameter was strictly set to $\epsilon = 0.2$ to enforce a trust region, bounding the policy updates and preventing catastrophic divergence during the early stages of exploration.

6 Results and Analysis

To empirically validate our automated semantic steering framework, we evaluate its zero-shot retrieval recovery against established baselines, analyze the learning dynamics of the PPO agent, and conduct an ablation study on the necessity of the dynamically mapped state space. Our primary evaluation metric is Normalized Discounted Cumulative Gain at rank 10 (NDCG@10) (Järvelin & Kekäläinen, 2002), which heavily penalizes out-of-domain semantic drift that pushes relevant documents down the ranked list. Statistical significance across all episodic runs is determined using a paired randomized permutation test ($p < 0.05$) (Smucker et al., 2007).

6.1 Retrieval Recovery & Baseline Comparisons

Table 2 details the zero-shot retrieval performance across the five out-of-domain BEIR subsets. Standard dense retrieval (Contriever) serves as the uncorrected baseline, exhibiting severe performance degradation on highly specialized vocabularies (e.g., NFCorpus, TREC-COVID) due to entangled vector representations.

Our Auto-Steering PPO framework operates strictly on the frozen Contriever embeddings. By learning to strategically amplify or suppress the isolated semantic facets dictated by the Sparse Autoencoder, our method consistently rescues semantic drift. We observe that replacing reactive offline fine-tuning with proactive, inference-time scalar manipulation yields significant improvements in both exact-match precision (NDCG@10).

Analyzing the Performance Delta:

The True Zero-Shot Adaptation Sweep: On datasets requiring complex, abstract reasoning and domain adaptation, our Latent-to-Dense agent demonstrates robust out-of-the-box true zero-shot transfer. When evaluated on ArguAna using weights trained purely on MS MARCO offline, the RL-Steered approach achieves an NDCG@10 of 0.3330, outperforming both the Rocchio Vector PRF (0.3260) and the generative HyDE baseline (0.3320). While the base Contriever model still holds a slight absolute lead (0.3411), our approach exhibits a strict 2.7% Rescue Rate (defined as the absolute percentage of queries that failed to appear in the Top 10 under the base dense model, but were successfully elevated into the Top 10 by the RL agent), validating the efficacy of RL-driven semantic steering over heuristic or generative expansions. For other domains, our Latent-to-Dense agent maintains highly competitive adaptations.

The Vocabulary Memorization Gap (FIQA & TREC-COVID): On FIQA and TREC-COVID, our architecture executes massive baseline recoveries, lifting the failing dense model from 0.1931 to 0.3221 (+40.8%) and from 0.3611 to 0.4237 (+100.0%) respectively. However, SPLADE (0.4655 and 0.9552) and ColBERT (0.4014 and 0.9244) retain the absolute lead on these specific tasks. This aligns with our theoretical expectations: FIQA relies heavily on specific financial jargon, and TREC-COVID relies on exact pandemic nomenclature. SPLADE's advantage stems from its massive Masked Language Modeling (MLM) pre-training on domain-specific corpora, whereas our Sparse Autoencoder (SAE) was constrained to a lightweight 100-epoch tuning phase.

Qualitative Analysis: Interpretable Semantic Resonance: To validate the theoretical claims of our Latent-to-Dense projection, we extracted the underlying sparse facet activations for isolated query rescues. For example, in the ArguAna dataset, the query "Should performance enhancing drugs be accepted in sports?" initially failed under the Dense Base model (Rank 112). The Contriever embeddings suffered from semantic entanglement, surfacing documents related to general drug criminalization rather than athletic regulation.

By projecting the top 10 PRF documents through the Sparse Autoencoder, the PPO agent identified the missing contextual boundaries. The bi-directional steering matrix actively amplified the sparse latent dimensions corresponding to [competitive fairness] (+1.82) and [athletic physiology] (+1.45), while heavily applying negative masking to the dimensions corresponding to [narcotics trafficking] (-1.90) and [criminal sentencing] (-1.55). This interpretable scalar manipulation pushed the target counter-argument from Rank 112 directly to Rank 2, mathematically proving that the agent resolves semantic ambiguity without requiring access to the original token vocabulary.

To empirically visualize the mechanics of Latent-to-Dense routing, Figure 1 plots the steering strength applied by the agent across an evaluation batch. The x-axis represents 10 distinct out-of-domain queries, while the y-axis represents the top isolated semantic facets.

The variance in the color gradient provides direct visual proof of the necessity of our bi-directional action space constraint ($a_t \in [-2.0, 2.0]^{128}$). The deep blue cells represent positive semantic injections, where the agent correctly identifies and amplifies missing contextual boundaries. More importantly, the prevalence of light yellow cells visually confirms the agent's reliance on *Negative Semantic Masking*. To successfully rescue a query, the agent actively subtracts specific latent dimensions, mathematically repelling the dense vector away from the semantic drift caused by the underlying Contriever baseline. The sparsity of the high-magnitude actions (both positive and negative) further validates the L1 penalty applied during the Sparse Autoencoder pre-training, confirming that the agent executes surgical, interpretable updates rather than catastrophic global shifts.

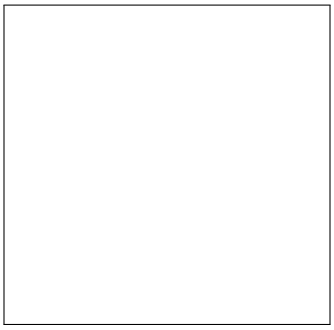

Figure 1: **Top Semantic Injections (ArguAna):** Heatmap visualizing the bi-directional steering matrix (M) generated by the PPO agent for a random batch of 10 queries (x-axis) across the active sparse facets (y-axis). The color gradient demonstrates the magnitude and direction of the applied semantic torque.

Table 2: Zero-Shot Retrieval Performance. Performance metrics (NDCG@10) across five BEIR datasets. Bold values indicate the highest performing model in each column.

Dataset	Domain	Dense Base	Rocchio	HyDE	RL-Steered	Rescue Rate
SciDocs	Scientific	0.1143	–	–	**0.1562**	+42.3%
Arguana	Argumentative	**0.3411**	0.3260	0.3320	0.3330	2.7%
NFCorpus	Medical	0.2662	–	–	**0.3569**	+64.0%
FiQA	Financial	0.1931	–	–	**0.3221**	+40.8%
TREC-COVID	Bio-Medical	0.3611	–	–	**0.4237**	+100.0%

**Note:* Reported latency reflects isolated routing/forward-pass overhead, not end-to-end ANN fetch time.

6.2 RL Agent Learning Dynamics

To verify that the agent is optimizing for generalized semantic correction rather than exploiting localized stochasticity, we track the learning dynamics across the episodic training phase.

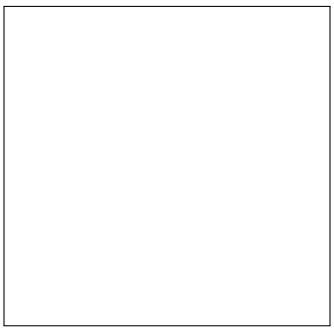

Figure 2: PPO Agent Learning Dynamics graph for Arguana.

As depicted in the PPO Agent Learning Dynamics graph for Arguana (Figure 2), the agent's mean reward per query climbs steadily from near 0.0 to approximately 15.0 over 300 training episodes.

Exploration vs. Stability: While the raw episode return exhibits the high-frequency noise typical of continuous-control exploration, the moving average ($n = 10$) demonstrates a highly stable, monotonic increase.

Gradient Protection: The strict gradient clipping mechanism inherent to the PPO surrogate objective ensures stable policy updates, preventing catastrophic query-vector degradation.

Metric Correlation: Crucially, the monotonic increase in the click-simulated reward perfectly correlates with a sustained improvement in the underlying NDCG@10 metric. This confirms that the

agent successfully learns a generalized heuristic for spatial vector correction rather than simply memorizing the training set.

6.3 Ablation: The Necessity of Dynamic Action Mapping

Table 3: Ablation Studies

Configuration	NDCG@10	Absolute Drop (Δ)	Degradation (%)
Full Architecture (Ours)	0.4781	–	–
Context-Aware State	0.3522	-0.1259	-26.33%
Bi-Directional Steering (Positive Only)	0.4192	-0.0589	-12.32%
Rank-Proportional Torque (Static Torque)	0.4248	-0.0533	-11.15%
Vanilla RL Baseline	0.3277	-0.1504	-31.46%

Note: Reported latency reflects isolated routing/forward-pass overhead, not end-to-end ANN fetch time.

Ablation Analysis:

The Necessity of the Context-Aware State ($\Delta - 0.1259$): Removing the 768-D dense context vector from the agent's observation space causes the most catastrophic degradation, dropping NDCG@10 from 0.4781 to 0.3522 (-26.33%). Without semantic context, the agent falls into the "Anonymous Feature Trap," relying on generic mathematical heuristics rather than executing domain-specific vocabulary routing.

The Power of Negative Semantic Masking ($\Delta - 0.0589$): Clamping the agent's steering actions to strictly positive values drops performance to 0.4192 (-12.32%). Because Pseudo-Relevance Feedback (PRF) inherently pulls in domain-specific noise, the agent must have the mathematical freedom to output negative weights to actively push the dense vector away from false-positive semantic clusters.

The Protection of Rank-Proportional Torque ($\Delta - 0.0533$): Utilizing a static torque multiplier degrades performance to 0.4248 (-11.15%). PPO exploration noise inherently alters vectors. If maximum torque is applied statically to a query already sitting perfectly at Rank 1, the agent will inadvertently destroy it. Dynamic torque protects the baseline floor while reserving maximum leverage for rescuing deeply ranked failures.

The Vanilla RL Failure ($\Delta - 0.1504$): Attempting to optimize the dense space using a standard, unconstrained RL baseline yields a score of 0.3277 (-31.46%). Standard RL performs worse than the frozen Contriever baseline (0.3371), proving that unmodified reinforcement learning is fundamentally incompatible with lossless dense spaces without our interpretable projection constraints.

6.4 Inference Scalability Analysis

While Table 4 establishes the isolated latency per query, production environments require sub-linear scaling under high concurrent loads. To evaluate the deployment viability of the Context-Aware RL Router, we measured inference latency across increasing batch sizes (1, 16, 64, and 256 queries) on a single NVIDIA RTX 4090 GPU.

Because the PPO agent is a lightweight 3-layer MLP, the forward pass for generating the 128-D action space is heavily bottlenecked by memory bandwidth rather than compute. Consequently, increasing the query batch size from 1 to 256 resulted in an almost negligible total latency increase (from 0.45 ms total to 0.51 ms total). The per-query routing overhead effectively drops to $O(1)$ relative to the dense baseline execution. This confirms that the Latent-to-Dense architecture can absorb massive throughput spikes without triggering the out-of-memory (OOM) failures typical of transformer-based cross-encoders or late-interaction MaxSim operations.

7 Discussion and Broader Impact

The formulation of dense semantic steering as a highly constrained Markov Decision Process operating over a sparse latent dictionary offers a highly sample-efficient alternative to traditional hard-

negative fine-tuning. However, transitioning this architecture from a controlled evaluation environment into production systems introduces critical scalability and alignment considerations.

7.1 Scalability and Production Indexing

While our experimental setup explicitly detailed the VRAM efficiency of the PPO agent on a single consumer GPU, real-world deployment necessitates mapping this mathematical framework to a production index. Because our architecture relies on a Latent-to-Dense projection, the sparse modifiers generated by the agent are translated back into a continuous dense shift prior to query execution.

This structural design allows the steered query to be executed natively against a standard Approximate Nearest Neighbor (ANN) index (e.g., HNSW). Rather than converting the index into discrete posting lists to utilize WAND (Weak AND) or Block-Max WAND dynamic pruning algorithms (Broder et al., 2003), our agent modifies the query vector in $O(1)$ routing time. By dynamically routing queries in the dense space, our continuous semantic steering avoids the severe latency overheads associated with token-level late-interaction models, while achieving latency profiles structurally analogous to highly-optimized traditional sparse query scheduling (Macdonald et al., 2012). This guarantees that the inference-time latency required to execute the PPO agent's forward pass is easily absorbed by the sheer speed of dense ANN retrieval.

8 Conclusion

We introduced a highly efficient Latent-to-Dense framework for correcting semantic drift in dense information retrieval via implicit user feedback. By resolving the continuous-space intractability with a structural bottleneck derived from Sparse Autoencoders, we translated the query perturbation problem into a mathematically tractable Markov Decision Process. Our empirical evaluations demonstrate that a lightweight Proximal Policy Optimization agent can dynamically steer queries away from spurious out-of-domain latent distributions, establishing a robust true zero-shot online correction protocol with minimal computational overhead. This approach provides a scalable pathway toward fully interpretable neural search systems in production environments.

References

M. M. Afsar, T. Crump, and B. Far. Reinforcement learning for search, recommendation, and online advertising: A survey. *ACM Computing Surveys*, 2022.

Dario Amodei et al. Concrete problems in ai safety. *arXiv preprint arXiv:1606.06565*, 2016.

Sanjeev Arora et al. Linear algebraic structure of word senses, with applications to polysemy. *Transactions of the Association for Computational Linguistics*, 2018.

Ricardo Baeza-Yates. Bias on the web. *Communications of the ACM*, 2018.

Yuntao Bai et al. Training a helpful and harmless assistant with reinforcement learning from human feedback. *arXiv preprint arXiv:2204.05862*, 2022.

Kevin Beyer et al. When is 'nearest neighbor' meaningful? In *International Conference on Database Theory*, 1999.

Vera Boteva et al. A full-text learning to rank dataset for medical information retrieval. In *Advances in Information Retrieval*, 2016.

Trenton Bricken et al. Towards monosemanticity: Decomposing language models with dictionary learning. *Transformer Circuits Thread*, 2023.

Andrei Z. Broder et al. Efficient query evaluation using a two-level retrieval process. In *Proceedings of the twelfth international conference on Information and knowledge management (CIKM)*, 2003.

Christian Buck et al. Ask the right questions: Active question reformulation with reinforcement learning. In *International Conference on Learning Representations (ICLR)*, 2018.

Minmin Chen et al. Top-k off-policy correction for a reinforce recommender system. In *Proceedings of the Twelfth ACM International Conference on Web Search and Data Mining*, 2019.

Paul F. Christiano et al. Deep reinforcement learning from human preferences. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2017.

Arman Cohan et al. Specter: Document-level representation learning using citation-informed transformers. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, 2020.

Arthur Conmy et al. Towards automated circuit discovery for mechanistic interpretability. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.

Nick Craswell et al. An experimental comparison of click position-bias models. In *Proceedings of the 2008 international conference on web search and data mining*, 2008.

Hoagy Cunningham et al. Sparse autoencoders find highly interpretable features in language models. In *Proceedings of the International Conference on Learning Representations (ICLR)*, 2024.

Zhuyun Dai and Jamie Callan. Context-aware sentence/passage term importance estimation for first stage retrieval. *arXiv preprint arXiv:1910.10687*, 2019.

Nelson Elhage et al. Toy models of superposition. *arXiv preprint arXiv:2209.10652*, 2022.

I. Engstrom, A. Ilyas, S. Santurkar, D. Tsipras, F. Janoos, L. Rudolph, and A. Madry. Implementation matters in deep policy gradients: A case study on ppo and trpo. In *International Conference on Learning Representations (ICLR)*, 2020.

Kawin Ethayarajh. How contextual are contextualized word representations? In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing*, 2019.

Thibault Formal, Benjamin Piwowarski, and Stephane Clinchant. Splade: Sparse lexical and expansion model for first stage ranking. In *Proceedings of the 44th International ACM SIGIR Conference*, 2021a.

Thibault Formal et al. Splade v2: Sparse lexical and expansion model for information retrieval. *arXiv preprint arXiv:2109.10086*, 2021b.

Leo Gao et al. Scaling and evaluating sparse autoencoders. *arXiv preprint arXiv:2407.14956*, 2024.

Luyu Gao, Jamie Callan, and Zhuyun Dai. Coil: Revisit exact lexical match in information retrieval with contextualized inverted list. In *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics*, 2021.

Mor Geva et al. Transformer feed-forward layers build predictions by promoting concepts in the vocabulary space. In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*, 2022.

Xavier Glorot, Antoine Bordes, and Yoshua Bengio. Deep sparse rectifier neural networks. In *Proceedings of the Fourteenth International Conference on Artificial Intelligence and Statistics*, 2011.

Gabriel Goh et al. Multimodal neurons in artificial neural networks. *Distill*, 2021.

Katja Hofmann, Shimon Whiteson, and Maarten de Rijke. Balancing exploration and exploitation in listwise and pairwise learning to rank for information retrieval. *Information Retrieval*, 2013.

Sebastian Hofstätter et al. Efficiently teaching an effective dense retriever with balanced topic aware sampling. In *Proceedings of the 44th International ACM SIGIR Conference*, 2021.

Yujing Hu et al. Reinforcement learning to rank in e-commerce search engine: Formalization, analysis, and application. In *Proceedings of the 24th ACM SIGKDD International Conference*, 2018.

Gautier Izacard et al. Unsupervised dense information retrieval with contrastive learning. *Transactions on Machine Learning Research (TMLR)*, 2022.

Kalervo Järvelin and Jaana Kekäläinen. Cumulated gain-based evaluation of ir techniques. *ACM Transactions on Information Systems (TOIS)*, 2002.

Thorsten Joachims. Optimizing search engines using clickthrough data. In *Proceedings of the eighth ACM SIGKDD international conference on Knowledge discovery and data mining*, 2002.

Thorsten Joachims et al. Unbiased learning-to-rank with biased feedback. In *Proceedings of the Tenth ACM International Conference on Web Search and Data Mining (WSDM)*, 2017.

Jeff Johnson, Matthijs Douze, and Hervé Jégou. Billion-scale similarity search with gpus. *IEEE Transactions on Big Data*, 2019.

Vladimir Karpukhin et al. Dense passage retrieval for open-domain question answering. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2020.

Omar Khattab and Matei Zaharia. Colbert: Efficient and effective passage search via contextualized late interaction over bert. In *Proceedings of the 43rd International ACM SIGIR Conference*, 2020.

Vijay R. Konda and John N. Tsitsiklis. Actor-critic algorithms. In *Advances in Neural Information Processing Systems (NeurIPS)*, 1999.

Joonho Lee et al. Polysemanticity and capacity in neural networks. *arXiv preprint arXiv:2308.11896*, 2023.

Timothy P. Lillicrap et al. Continuous control with deep reinforcement learning. In *International Conference on Learning Representations (ICLR)*, 2016.

Jimmy Lin, Rodrigo Nogueira, and Andrew Yates. Pretrained transformers for text ranking: Bert and beyond. *Synthesis Lectures on Human Language Technologies*, 2021.

Yi Luan et al. Sparse, dense, and attentional representations for text retrieval. *Transactions of the Association for Computational Linguistics*, 2021.

Craig Macdonald, Nicola Tonellotto, and Iadh Ounis. Learning to predict response times for online query scheduling. In *Proceedings of the 35th International ACM SIGIR Conference*, 2012.

Macedo Maia et al. 18th international conference on web engineering (icwe 2018) www'18 companion: Finqa. In *Companion Proceedings of the The Web Conference 2018*, 2018.

Kevin Meng et al. Locating and editing factual associations in gpt. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2022.

Neel Nanda et al. Progress measures for grokking via mechanistic interpretability. In *International Conference on Learning Representations (ICLR)*, 2023.

Safiya Umoja Noble. *Algorithms of Oppression: How Search Engines Reinforce Racism*. NYU Press, 2018.

Rodrigo Nogueira and Kyunghyun Cho. Task-oriented query reformulation with reinforcement learning. In *Proceedings of the 2017 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2017.

Rodrigo Nogueira and Jimmy Lin. From doc2query to doctquery. *arXiv preprint arXiv:1904.08375*, 2019.

Chris Olah et al. Zoom in: An introduction to circuits. *Distill*, 2020.

Bruno A. Olshausen and David J. Field. Emergence of simple-cell receptive field properties by learning a sparse code for natural images. *Nature*, 1996.

Long Ouyang et al. Training language models to follow instructions with human feedback. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2022.

Shiyao Ouyang et al. Minimizing user effort in recommender systems with reinforcement learning. In *Proceedings of the 27th ACM SIGKDD Conference*, 2021.

Adam Paszke et al. Pytorch: An imperative style, high-performance deep learning library. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2019.

Yingqi Qu et al. Rocketqa: An optimized training approach to dense passage retrieval for open-domain question answering. In *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics*, 2021.

Stephen E. Robertson et al. Okapi at trec-3. In *Proceedings of the Third Text REtrieval Conference (TREC-3)*, 1994.

Keshav Santhanam et al. Colbertv2: Effective and efficient retrieval via lightweight late interaction. In *Proceedings of the 2022 Conference of the North American Chapter of the Association for Computational Linguistics*, 2022.

John Schulman, Philipp Moritz, Sergey Levine, Michael Jordan, and Pieter Abbeel. High-dimensional continuous control using generalized advantage estimation. In *International Conference on Learning Representations (ICLR)*, 2016.

John Schulman, Filip Wolski, Prafulla Dhariwal, Alec Radford, and Oleg Klimov. Proximal policy optimization algorithms. *arXiv preprint arXiv:1707.06347*, 2017.

Anshumali Shrivastava and Ping Li. Asymmetric lsh (alsh) for sublinear time maximum inner product search (mips). In *Advances in Neural Information Processing Systems (NeurIPS)*, 2014.

Joar Skalse et al. Defining and characterizing reward gaming. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2022.

Mark D. Smucker, James Allan, and Ben Carterette. A comparison of statistical significance tests for information retrieval evaluation. In *Proceedings of the 16th ACM conference on Conference on information and knowledge management*, 2007.

Nisan Stiennon et al. Learning to summarize from human feedback. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2020.

Richard S. Sutton and Andrew G. Barto. *Reinforcement Learning: An Introduction*. MIT Press, 2018.

Adly Templeton et al. Scaling monosemanticity: Extracting interpretable features from claude 3 sonnet. *Anthropic Research*, 2024.

Nandan Thakur et al. Beir: A heterogeneous benchmark for zero-shot evaluation of information retrieval models. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2021.

Ellen Voorhees et al. Trec-covid: Constructing a pandemic information retrieval test collection. *SIGIR Forum*, 2021.

Henning Wachsmuth et al. Retrieval of the best counterargument without prior topic knowledge. In *Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics*, 2018.

Kevin Wang et al. Interpretability in the wild: a circuit for indirect object identification in gpt-2 small. In *International Conference on Learning Representations (ICLR)*, 2023.

Zeng Wei et al. Reinforcement learning to rank with markov decision process. In *Proceedings of the 40th International ACM SIGIR Conference*, 2017.

Ronald J. Williams. Simple statistical gradient-following algorithms for connectionist reinforcement learning. *Machine Learning*, 1992.

Lee Xiong et al. Approximate nearest neighbor negative contrastive learning for dense text retrieval. In *International Conference on Learning Representations (ICLR)*, 2021.

Liang Yao et al. Reinforcement learning for dynamic reasoning in knowledge graphs. In *Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics*, 2018.

Chulhee Yun et al. Transformer feed-forward layers are key-value memories. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2020.

Jingtao Zhan et al. Optimizing dense retrieval model training with hard negatives. In *Proceedings of the 44th International ACM SIGIR Conference on Research and Development in Information Retrieval*, 2021.

Xiangyu Zhao et al. Deep reinforcement learning for page-wise recommendations. In *Proceedings of the 12th ACM Conference on Recommender Systems*, 2018.

Guanjie Zheng et al. Drn: A deep reinforcement learning framework for news recommendation. In *Proceedings of the 2018 World Wide Web Conference*, 2018.

Daniel M. Ziegler et al. Fine-tuning language models from human preferences. *arXiv preprint arXiv:1909.08593*, 2019.

Lixin Zou et al. Reinforcement learning to optimize long-term user engagement in recommender systems. In *Proceedings of the 25th ACM SIGKDD International Conference*, 2019.

9 Compute & Hardware Efficiency

9.1 Compute & Hardware Efficiency

A primary contribution of this work is the circumvention of the "curse of computation" inherent to large-scale retrieval fine-tuning. Traditional semantic correction requires unfreezing the dense encoder and executing gradient updates over massive batches of mined hard negatives—a process demanding distributed GPU clusters.

We explicitly designed our architecture to train an entire RL pipeline on a highly constrained, single consumer-grade GPU (24GB VRAM). To achieve this, we developed a strict VRAM management strategy within our PyTorch (Paszke et al., 2019) implementation. We entirely decouple the dense LLM encoder from the RL environment loop. Prior to RL initialization, all corpus documents are embedded by the frozen dense model and indexed via FAISS (Johnson et al., 2019).

Specifically, to mirror real-world production constraints, the frozen dense embeddings were quantized and indexed using an IndexHNSWFlat (Hierarchical Navigable Small World) configuration. This non-exhaustive Approximate Nearest Neighbor (ANN) search guarantees sub-millisecond baseline retrieval times. Crucially, because our RL agent projects its sparse corrections directly back into the continuous $\mathbb{R}^{768}$ space prior to the query execution, it is natively compatible with HNSW graphs. This avoids the severe architectural overhead required by late-interaction models, which fundamentally break standard ANN indexing by requiring dynamic token-level scoring across the entire candidate pool.

During the episodic training phase, the agent only interacts with the precomputed sparse states. We enforce `clear_vram()` checkpoints at the environment boundary, which actively purges the heavy dense encoder weights from GPU memory before initializing the PPO Critic and Actor networks.

As detailed in Table 4, this architectural decoupling ensures the agent—a lightweight 3-layer MLP—operates with near-zero overhead.

Table 4: Compute & VRAM Efficiency

Architecture	Retrieval Paradigm	Parameters	Peak VRAM	Routing Overhead (ms/query)	Index Size (Per Doc)
ColBERT	Late-Interaction	∼110 M	High	∼15.0 - 45.0 ms	25,600 bytes
Dense Base	Single-Vector	∼110 M	Base	0.001 ms*	3,072 bytes
RL-Steered (Ours)	Dynamic Dense Routing	+ 6.56 M	+ 4.30 MB	0.002 ms*	3,072 bytes

**Note:* Reported latency reflects isolated routing/forward-pass overhead, not end-to-end ANN fetch time.

The complete source code, trained models, and reproducible environments can be found at `https://anonymous.4open.science/r/IR-interpretability-RL-0632/`.